MAST32001 Computational Statistics

Course Assignment Task 2: Bayesian mixture modelling with MCMC

Kaiwei Fu Lingyu Wang Sami Pönkä Zhi Zhou

21 October 2026

Task 2: Bayesian mixture modelling with MCMC

Part A: Latent-variable representation and the Gibbs sampler

Model specification and priors

Each observation comes from one of two normal components, but its source is unknown. The observations are independent, so we do not need the block-resampling machinery designed for dependent time series (**?**). For independent observations $y_1, \dots, y_n$, introduce latent allocations $Z_i \in \{1, 2\}$, where $\Pr(Z_i = 1) = \pi$ and $\Pr(Z_i = 2) = 1 - \pi$. Conditional on its allocation,

$$Y_i \mid Z_i = k, \mu_1, \mu_2, \sigma^2 \sim \mathcal{N}(\mu_k, \sigma^2), \quad k = 1, 2.$$

Integrating out Z_i gives the familiar mixture density:

$$p(y_i \mid \mu_1, \mu_2, \pi, \sigma^2) = \pi\varphi(y_i; \mu_1, \sigma^2) + (1 - \pi)\varphi(y_i; \mu_2, \sigma^2),$$

where $\varphi(\cdot; \mu, \sigma^2)$ denotes the normal density.

We use the independent, proper priors

$$\mu_1, \mu_2 \overset{\text{iid}}{\sim} \mathcal{N}(m_0, s_0^2), \quad \pi \sim \text{Beta}(a, b), \quad \sigma^2 \sim \text{InvGamma}(\alpha_0, \beta_0),$$

with $m_0 = 0$, $s_0 = 10$, $a = b = 1$, and $\alpha_0 = 2$, $\beta_0 = 1$. The mean prior is broad on the scales. The weight prior is uniform, and the two mean priors are identical, so the model treats the two labels symmetrically. The variance prior has mean 1 and infinite variance. These choices keep every update conjugate. They do not make the priors irrelevant: a component with few observations is still sensitive to them.

Our inverse-gamma convention is

$$p(v) \propto v^{-\alpha_0 - 1} \exp(-\beta_0 / v), \quad v > 0,$$

so $1/v$ is gamma with shape α_0 and rate β_0. This convention fixed the rate used in the R sampler.

To derive the updates, let $n_k = \sum_i \mathbb{I}(Z_i = k)$ and $S_k = \sum_{i:Z_i=k} y_i$. Given the allocations, each observation contributes to only one normal component. The likelihood is

$$L(\theta, z; y) \propto \pi^{n_1}(1 - \pi)^{n_2}(\sigma^2)^{-n/2} \exp\left[-\frac{1}{2\sigma^2}\sum_i (y_i - \mu_{Z_i})^2\right].$$

Multiplying by the independent priors gives $p(\theta, z \mid y) \propto L(\theta, z; y)p(\mu_1)p(\mu_2)p(\pi)p(\sigma^2)$. This is the basis of the four conditional updates. Each update keeps only the terms involving its own unknown.

Full conditional distribution

The complete-data form lets us update one block at a time. Each distribution below conditions on the current values of the other unknown.

1. **Allocations.** Given the current parameters, Bayes' rule gives the probability that observation i belongs to component 1:

$$\Pr(Z_i = 1 \mid y_i, \theta) = \frac{\pi\varphi(y_i; \mu_1, \sigma^2)}{\pi\varphi(y_i; \mu_1, \sigma^2) + (1-\pi)\varphi(y_i; \mu_2, \sigma^2)} = p_i.$$

   We sample $\mathbb{I}(Z_i = 1) \sim \text{Bernoulli}(p_i)$ independently across observations. To avoid numerical underflow, it first calculates $l_{i1} = \log \pi + \log \varphi(y_i; \mu_1, \sigma^2)$, $l_{i2} = \log(1-\pi) + \log \varphi(y_i; \mu_2, \sigma^2)$, then uses $p_i = \text{logit}^{-1}(l_{i1} - l_{i2})$.

2. **Component means.** Once observations have been allocated, each component has a normal likelihood. For $k = 1, 2$, combining it with the mean prior gives

$$p(\mu_k \mid -) \propto \exp\left[-\frac{(\mu_k - m_0)^2}{2s_0^2} - \frac{1}{2\sigma^2}\sum_{i:Z_i=k}(y_i - \mu_k)^2\right].$$

   Completing the square gives

$$V_k = \left(s_0^{-2} + \frac{n_k}{\sigma^2}\right)^{-1}, \quad M_k = V_k\left(\frac{m_0}{s_0^2} + \frac{S_k}{\sigma^2}\right), \quad \mu_k \mid - \sim \mathcal{N}(M_k, V_k).$$

   If $n_k = 0$, then $S_k = 0$ and the update reduces to the prior. We will use this fact later, because it explains the large mean excursions in the overlap run.

3. **Mixing proportion.** The sampled counts contribute $\pi^{n_1}(1-\pi)^{n_2}$ to the likelihood. Combining this with the beta prior gives $\pi \mid - \sim \text{Beta}(a + n_1, b + n_2)$.

4. **Variance.** Finally, pool the residuals from both components as $\text{SSE} = \sum_i (y_i - \mu_{Z_i})^2$. Combining likelihood and prior gives

$$p(\sigma^2 \mid -) \propto (\sigma^2)^{-(\alpha_0 + n/2) - 1} \exp\left[-\frac{\beta_0 + \text{SSE}/2}{\sigma^2}\right],$$

   and hence $\sigma^2 \mid - \sim \text{InvGamma}\left(\alpha_0 + \frac{n}{2}, \beta_0 + \frac{\text{SSE}}{2}\right)$. The mean priors do not involve σ^2, so they add no extra term to this conditional shape.

Turning the conditionals into R code

With the four conditionals in hand, the sampler writes itself. It cycles through the updates, always conditioning on the most recent values. The R code is ‘rgibbs$_m ix()‘from‘R/task2.R‘.Ateveryiteration, wekeep(\mu$ and n_1, the number of observations allocated to component 1. That is enough for the posterior inference and for the diagnostics discussed in Parts B and C. The allocation step works on the log scale to avoid underflow, and the sampler checks its input and every sampled state.

Starting values matter here, because the posterior has two symmetric modes. If two chains never agree, we cannot trust the answer. So for each setting we start two chains: one with the labels swapped, the other with the means spread widely. Then we check that both converge to the same target. The exact starting schemes are listed in the simulation design (Part B).

How long to run is a practical question, and the answer differs by setting. The separated and the faithful setting run for 20,000 iterations with 5,000 burn-in. The overlapping setting runs for 60,000 iterations with 15,000 burn-in, because its posterior is harder to explore and mixes more slowly. Part B checks whether these choices are enough, using trace plots, the effective sample size (ESS), and the Gelman-Rubin $\hat{R}$ statistic.

Part B: Simulation study and computational behaviour

Simulation design

The first test for a new sampler is a self-check. We generate data from known parameters, run the sampler, and see whether the posterior finds its way back to the truth. We do this with $n = 500$ observations and two settings that differ in one important way: how far apart the components are.

1. Clearly separated components: $\mu_1 = -2$, $\mu_2 = 2$. The densities barely overlap, the two modes are easy to tell apart, and the model is close to identifiable.

2. Substantially overlapping components: $\mu_1 = -0.5$, $\mu_2 = 0.5$. The component densities overlap heavily, and a large fraction of the mass is ambiguous between the two components.

The two settings are not equally hard. When the components are far apart, the likelihood has two well-separated modes and the posterior concentrates near the true parameters. When they overlap, the likelihood is flat between the modes. The posterior becomes multimodal and symmetric under the label permutation $(\mu_1, \mu_2, \pi) \leftrightarrow (\mu_2, \mu_1, 1 - \pi)$. The chain can jump between the two mirror-image modes, and that jump is what we call label switching. This symmetry is the main obstacle we have to handle when we summarize the output.

For each setting we run two chains from different starting values: one with swapped labels and one with widely spread means. Table 1 lists the exact values and seeds. The two starts differ in mean positions, label orientation, mixing proportion, and variance. If two chains that start this far apart still agree, the agreement is not an accident.

Table 1: Initial values and seeds for the two chains in each setting.

Setting	Chain	Seed	μ_1	μ_2	π	σ^2
Separated	swap	101	2.000	-2.000	0.9	1.000
Separated	spread	102	-4.000	4.000	0.5	4.000
Overlap	swap	201	2.000	-2.000	0.9	1.000
Overlap	spread	202	-4.000	4.000	0.5	4.000
Faithful	swap	301	4.533	2.003	0.8	1.303
Faithful	spread	302	0.459	6.241	0.5	2.605

Recovery of the data-generating mixture

Once we have the chains, the first question is whether the sampler finds the mixture that generated the data. The posterior does not care about the order of the labels. So before we compare anything, we impose the order restriction $\mu_{(1)} < \mu_{(2)}$, relabel every draw accordingly, and pool the two chains. Table 2 and Table 3 then compare the posterior median and the 95% credible interval of each parameter with its true value.

Table 2: Separated components: posterior summaries vs truth (order-restricted, two chains pooled).

Parameter	Truth	Median	lo	hi
μ_1 (lower)	-2.0	-1.98	-2.10	-1.86
μ_2 (upper)	2.0	2.04	1.91	2.18
π	0.5	0.54	0.49	0.58
σ^2	1.0	0.91	0.80	1.05

In the separated setting the answer is clean. The sampler recovers the generating mixture almost exactly. The posterior medians of μ_1 and μ_2 are -1.98 and 2.04, close to the true values

Table 3: Overlapping components: posterior summaries vs truth (order-restricted, two chains pooled).

Parameter	Truth	Median	lo	hi
μ_1 (lower)	-0.5	-0.40	-2.58	-0.12
μ_2 (upper)	0.5	0.08	-0.20	2.25
π	0.5	0.49	0.01	0.99
σ^2	1.0	1.14	0.88	1.34

-2 and 2. The 95% credible intervals are narrow and contain the truth. The same holds for π (median 0.54) and σ^2 (median 0.91).

Overlap makes the answer noisier, but not wrong. After relabeling, the posterior still centers near the truth. The median of μ_1 is -0.40 (true -0.5), the median of μ_2 is 0.08 (true 0.5), and the median of π is 0.49. The credible intervals are much wider, because the likelihood is flat between the modes. So once the label ambiguity is removed, the sampler recovers the data-generating mixture in both settings. The price of overlap is precision.

Before pooling, we also ask whether the answer depends on where we started. Table 4 compares the two chains separately. The per-chain ordered medians agree closely in every setting. The separated and faithful chains in particular agree to the second decimal, while they remain in opposite raw label orientations. So the agreement is not an artefact of identical labeling.

Table 4: Per-chain posterior medians after relabeling (order restriction $\mu_{(1)} < \mu_{(2)}$) and label-switch counts.

Setting	Chain	μ_{lo} median	μ_{hi} median	π_{lo} median	Switches
Separated	swap	-1.979	2.044	0.537	0
Separated	spread	-1.982	2.044	0.536	0
Overlap	swap	-0.411	0.082	0.475	2984
Overlap	spread	-0.395	0.073	0.498	3442
Faithful	swap	2.048	4.297	0.361	0
Faithful	spread	2.049	4.297	0.361	0

The trace plots make the mechanism visible. In the separated setting the chains behave as a well-mixing sampler should. The draws oscillate around stable levels, and the labels never swap. In the overlapping setting the raw traces wander between the two equivalent modes. Every large jump in the overlap panels is a label switch that swaps the roles of μ_1 and μ_2.

MCMC diagnostics and computational behaviour

Good medians are not enough on their own. We also want to know that the chains really converged. We measure mixing and convergence on the order-restricted chains with the effective sample size (ESS) and the maximum Gelman-Rubin $\hat{R}$ statistic. Table 5 summarizes the computational behavior of each setting. We screen the chains with all $\hat{R} \leq 1.01$ and all bulk and tail ESS ≥ 400. Every setting passes.

The separated setting behaves as advertised. The minimum bulk ESS is 20,826 and $\hat{R}$ reaches at most 1.0001. The overlapping setting is the more interesting one. Its raw (unrelabeled) chains are dominated by label switching. The raw ESS collapses and $\hat{R}$ rises far above 1. After the order restriction the chains still mix slowly. The minimum bulk ESS is 2,878 and maximum $\hat{R}$ is 1.0067, but they have been run long enough for the diagnostic screen to pass.

Effect of overlap: label switching and latent allocations

Label switching. To see how often a chain changes its mind about which component is which, we counted, for each chain, how many times the ordering of the means changed between consecutive

Table 5: Computational behaviour per setting: chain length, maximum Gelman-Rubin statistic, minimum bulk and tail ESS, total sampler time, and the fraction of ambiguous latent allocations (posterior probability in 0.2–0.8).

Metric	Separated	Overlap	Faithful
Iterations	20000.0	60000.0	20000.0
Burn-in	5000.0	15000.0	5000.0
Retained draws	15000.0	45000.0	15000.0
Maximum ordered $\hat{R}$	1.0001	1.0067	1.0001
Minimum bulk ESS	20826.0	2878.0	27119.0
Minimum tail ESS	25182.0	1756.0	28254.0
Sampler time (seconds)	5.92	20.93	2.53
Ambiguous allocations (%)	2.00	100.00	0.40

Table 6: Per-parameter diagnostics of the order-restricted pooled chains: posterior median, Gelman–Rubin statistic, bulk and tail ESS, and MCSE of the posterior mean.

Metric	Separated	Overlap	Faithful
— in the works —			

draws. In the separated setting the labels never switched (0 switches across both chains). In the overlapping setting the combined count was 6,426 over the same post-burn period. Under overlap, naive posterior means of the raw draws are meaningless. Both component means are averaged toward zero. So we need the order restriction $\mu_{(1)} < \mu_{(2)}$ before summarizing the component-specific parameters.

Label-invariant posterior quantities. Which summaries actually suffer from label switching? Quantities that are symmetric under the permutation $(\mu_1, \mu_2, \pi) \leftrightarrow (\mu_2, \mu_1, 1-\pi)$ stay meaningful when computed directly from the raw chains. These include the posterior predictive density of a future observation, the marginal posterior of σ^2, the posterior distribution of the latent allocation probabilities, and any clustering-type summary derived from them.

Latent allocations. Overlap also shows up in how sure the posterior is about the latent allocations. For each observation we computed its posterior probability of belonging to the lower-mean component, averaged over the pooled relabeled draws. Under separation these probabilities concentrate near 0 or 1. Only 2% of the observations are ambiguous, with probability between 0.2 and 0.8. Under overlap, 100% of the observations fall in this ambiguous range.

Empty components and weak identification

An empty component means that its sampled allocation count was zero in a retained sweep. An extreme weight is a mixing proportion below 0.01 or above 0.99. Table 7 reports how often each happens in each setting. Empty components are a known issue in Bayesian mixture estimation (**?**). These draws matter. An empty component draws its mean from the broad normal prior. With no observations to constrain it, the mean is free to wander until the next allocation.

Table 7: Fraction of retained sweeps with an empty component or an extreme mixing weight.

Setting	Empty-component draws (%)	Extreme weight draws (%)
Separated	0.00	0.00
Overlap	3.56	7.33
Faithful	0.00	0.00

Part C: Label switching and identifiability

The symmetry of the likelihood
Part B showed label switching happening in practice. Here we look at why it happens and what can be done about it. The reason is a symmetry of the model. The mixture density is invariant under permutation of the component labels:

$$p(y_i \mid \mu_1, \mu_2, \sigma^2, \pi) = p(y_i \mid \mu_2, \mu_1, \sigma^2, 1-\pi).$$

The same invariance holds for the posterior, because the prior is symmetric in μ_1 and μ_2. The parameter vector (μ_1, μ_2, π) is then identifiable only up to this permutation. The posterior has two equivalent modes. A chain that can cross the low-probability region between them alternates between the two. Label switching is a feature of the model itself, not an artefact of our implementation.

The new content:

Identifiability restrictions and post-processing
Two remedies are available. The first is a constrained prior, for example $\mu_1 < \mu_2$ a priori. It makes the model identifiable by construction, but it changes the target posterior and complicates the sampler. The second, which we use here, is to run the sampler unconstrained and relabel the draws afterwards. We use the order restriction $\mu_{(1)} < \mu_{(2)}$. At each iteration we assign the smaller and the larger of (μ_1, μ_2) to the two components and flip π accordingly ($\pi \mapsto 1-\pi$ when $\mu_1 > \mu_2$).
Critique of Simple Order Restrictions ($\mu_{(1)} < \mu_{(2)}$)
While post-processing with $\mu_{\text{lo}} = \min(\mu_1, \mu_2)$ and $\mu_{\text{hi}} = \max(\mu_1, \mu_2)$ resolves label assignment, it introduces two distinct artifacts under heavy component overlap:

1. **Order-statistic truncation (repulsion bias):** Because marginal posterior distributions overlap across zero, folding the parameter space across $\mu_1 = \mu_2$ truncates the tails, artificially repelling posterior medians away from each other ($\mathbb{E}[\min(\mu_1, \mu_2)] \leq \min(\mathbb{E}[\mu_1], \mathbb{E}[\mu_2])$).

2. **Prior-tail contamination from empty components:** In 3.56% of overlapping sweeps, an empty component ($n_k = 0$) causes μ_k to update directly from the diffuse prior $\mathcal{N}(0, 10^2)$. The order statistic systematically channels extreme negative prior draws into μ_{lo} and positive draws into μ_{hi}, generating skewed credible intervals ($[-2.58, -0.12]$ and $[-0.20, 2.25]$). Conditioning on non-empty occupancy ($\min(n_1, n_2) \geq 5$) successfully eliminates these prior-tail excursions.

Part D: Application to the Old Faithful geyser data

So far everything has been simulated. To see whether the sampler works in practice, we fitted the two-component mixture to the eruption durations of the Old Faithful geyser ($n = 272$, from the datasets package). The data are strongly bimodal. They correspond to short and long eruption regimes, and we analyse them on the original scale in minutes. We run two chains from setting-adapted starting values. One start swaps the labels and uses the 20% and 80% quantiles of the durations. The other spreads the means beyond the observed range. Each chain runs for 20,000 iterations with 5,000 burn-in.

The posterior places the two components at 2.05 and 4.30 minutes. The proportion of short eruptions is 0.36, and the residual standard deviation is 0.37 minutes. This agrees well with the known two-regime structure of the geyser. On the original scale the components are well separated. The chains never switch labels (0 switches), and the minimum bulk ESS is 27,119. The posterior summaries are reliable.

Table 8: Old Faithful eruptions: posterior summaries (order-restricted, two chains pooled).

Parameter	Median	lo	hi
μ (short)	2.05	1.97	2.12
μ (long)	4.30	4.24	4.35
π (short)	0.36	0.31	0.42
σ^2	0.14	0.12	0.17

Posterior Predictive Check and Heteroscedastic Extension
Visual inspection of the fitted density indicates a localized discrepancy: the observed short-duration peak is sharper and taller than the fitted lower component, which is broader and lower in that region. The common-variance assumption $\sigma_1^2 = \sigma_2^2$ forces a compromise between tightly clustered short eruptions and more dispersed long eruptions.

To evaluate this quantitatively, we conduct a posterior predictive check (**?**). Using regime-specific dispersion statistics ($T_{\text{ratio}} = \text{SD}_{\text{long}}/\text{SD}_{\text{short}}$), the observed spread ratio is 1.540 ($SD_{\text{short}} = 0.267, SD_{\text{long}} = 0.411$). Across replicated datasets generated under the shared-variance model, the posterior predictive p-value is $p_B < 0.001$, confirming the inadequacy of the homoscedastic assumption.

Extending the model to allow component-specific variances ($\sigma_1^2 \neq \sigma_2^2$) via independent Inverse-Gamma updates yields $\sigma_{\text{short}} \approx 0.295$ min and $\sigma_{\text{long}} \approx 0.412$ min. This heteroscedastic formulation captures the sharp short-eruption peak while avoiding likelihood degeneracy.

Conclusions

Our Gibbs sampler successfully recovers parameters across separated and overlapping regimes. While label switching and empty components require careful post-processing and diagnostic screening, the Bayesian mixture framework provides a robust description of bimodal geophysical processes.

References